

# ARTIFICIAL INTELLIGENCE

INFORMED SEARCH      *AIMA* CH. 3.5 AND 3.6

*Lecture 4*

# Outline

- ◇ Greedy search
- ◇ A\* search
- ◇ Admissible heuristic
- ◇ Optimality of A\* search
- ◇ Consistent heuristic
- ◇ Relaxed problems

## Revision: Tree search

```
function TREE-SEARCH(problem, frontier) returns a solution, or failure
  INSERT(ROOT-NODE(problem.INITIAL-STATE), frontier)
  while not EMPTY?(frontier) do
    node ← REMOVE(frontier)
    if problem.GOAL-TEST applied to node.STATE succeeds return node
    for each action in problem.ACTIONS(node.STATE) do
      INSERT(CHILD-NODE(problem, node, action), frontier)
  return failure
```

A strategy is defined by choosing the order of node expansion

## Best-first search

What if we have problem-specific knowledge beyond the problem definition?

Idea: use an evaluation function  $f(n)$  for each node  
– estimate of “desirability”

⇒ Expand most desirable unexpanded node

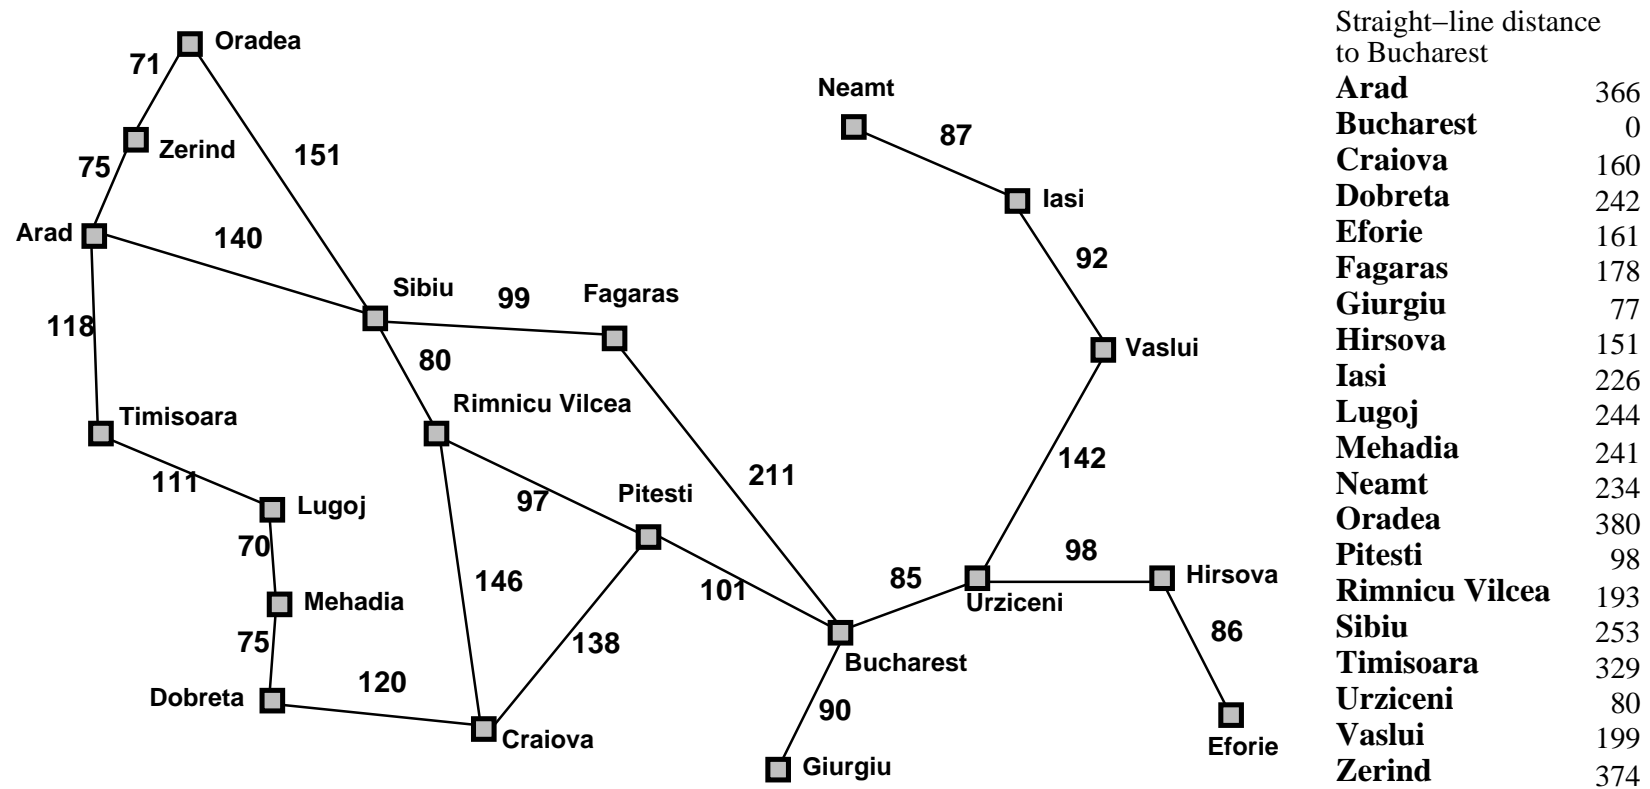
Implementation: *frontier* is a queue sorted in decreasing order of desirability

Special cases:

- greedy search
- A\* search

⇒ Obtained by different desirability notions

# Romania with step costs in km



## Greedy search

Evaluation function  $h(n)$  (**h**euristic)

= estimated cost of getting from node  $n$  to the closest goal

E.g.,  $h_{\text{SLD}}(n)$  = straight-line distance from  $n$  to Bucharest

Typically,  $h(n)$  depends only on the state of  $n$ , and not on its parents


– e.g., the straight-line distance from  $n$  to Bucharest does not depend on how we got to  $n$

Greedy search expands the node that **appears** to be closest to goal


◇ Resembles DFS in the way it prefers to follow a single path.

◇ Can go down an infinite path if we don't detect repeated states.

# Greedy search example

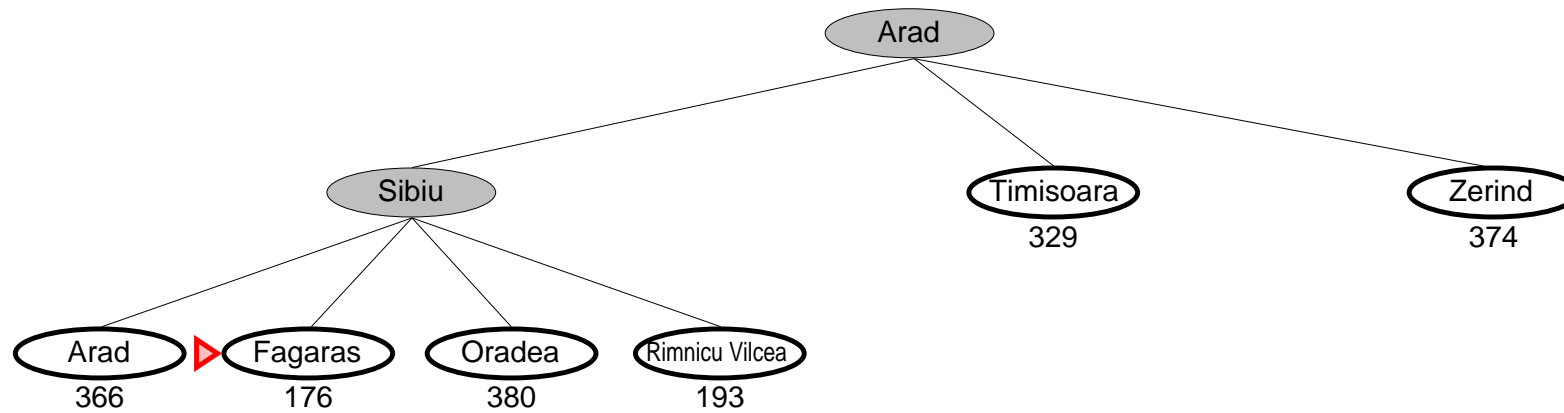


# Greedy search example

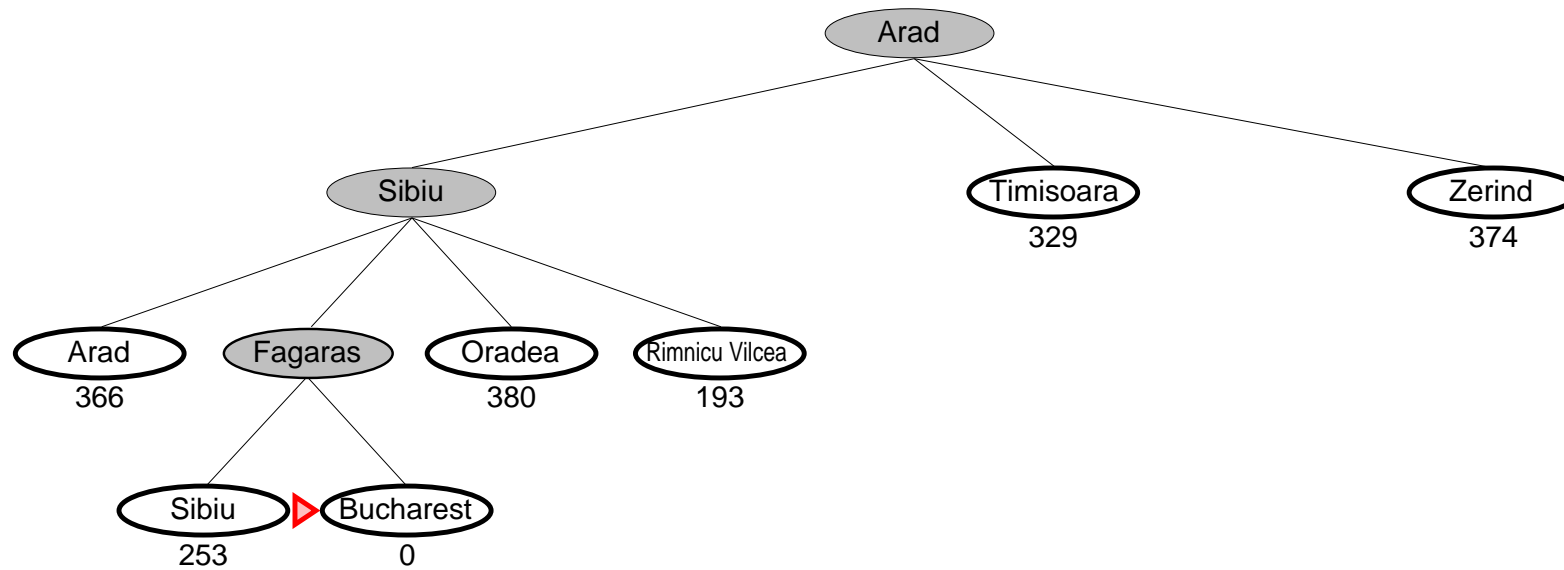




# Greedy search example



# Greedy search example



## Properties of greedy search

Complete?? No: can get stuck in loops

- e.g., with Oradea as goal:

Iasi  $\rightarrow$  Neamt  $\rightarrow$  Iasi  $\rightarrow$  Neamt  $\rightarrow \dots$

- Complete in finite state spaces with repeated-state checking

Time??  $O(b^m)$ , but a good heuristic can give dramatic improvement

Space??  $O(b^m)$  — keeps all nodes in memory

Optimal?? No

## A\* search

Idea: avoid expanding paths that are already expensive

A\* evaluation function:  $f(n) = g(n) + h(n)$

$g(n)$  = cost of the path from the start node to  $n$

– depends only on path cost, not on the heuristic

$h(n)$  = estimated cost of getting from node  $n$  to the closest goal

$f(n)$  = estimated total cost of a path through  $n$  to a goal

## Admissible heuristic

Let  $h^*(n)$  be the **true** cost of getting from  $n$  to the closest goal

Heuristic  $h(n)$  is **admissible** if  $0 \leq h(n) \leq h^*(n)$  for each node  $n$

(Note: these two conditions imply  $h(n_G) = 0$  for each goal node  $n_G$ )


$\Rightarrow$  I.e.,  $h(n)$  is admissible if it never overestimates the cost of getting from  $n$  to the closest goal

E.g.,  $h_{\text{SLD}}(n)$  never overestimates the actual road distance


**Theorem:** A\* **tree search** with an admissible heuristic is optimal

$\Rightarrow$  A\* tree search with admissible heuristic is used throughout AI!

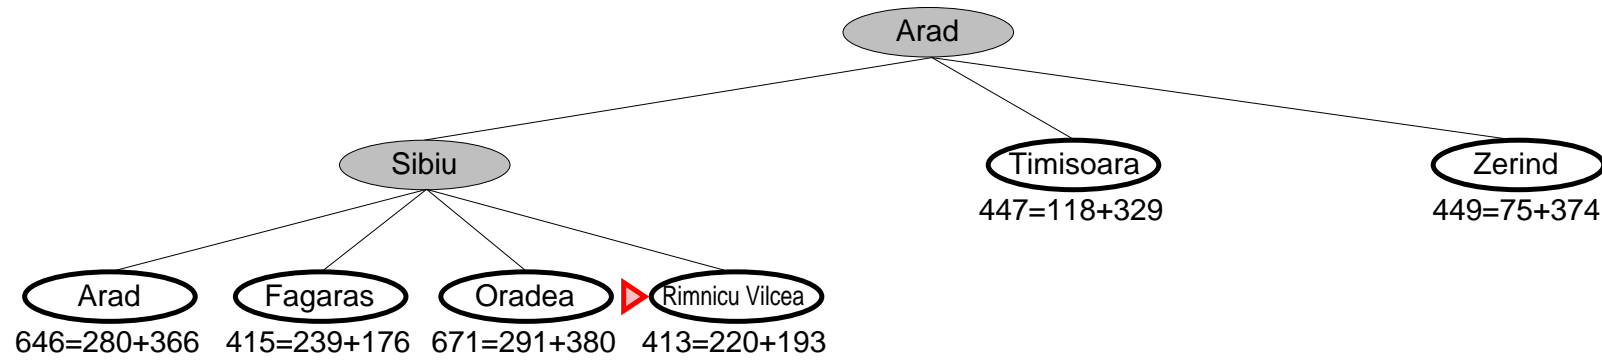
# A\* tree search example



# A\* tree search example

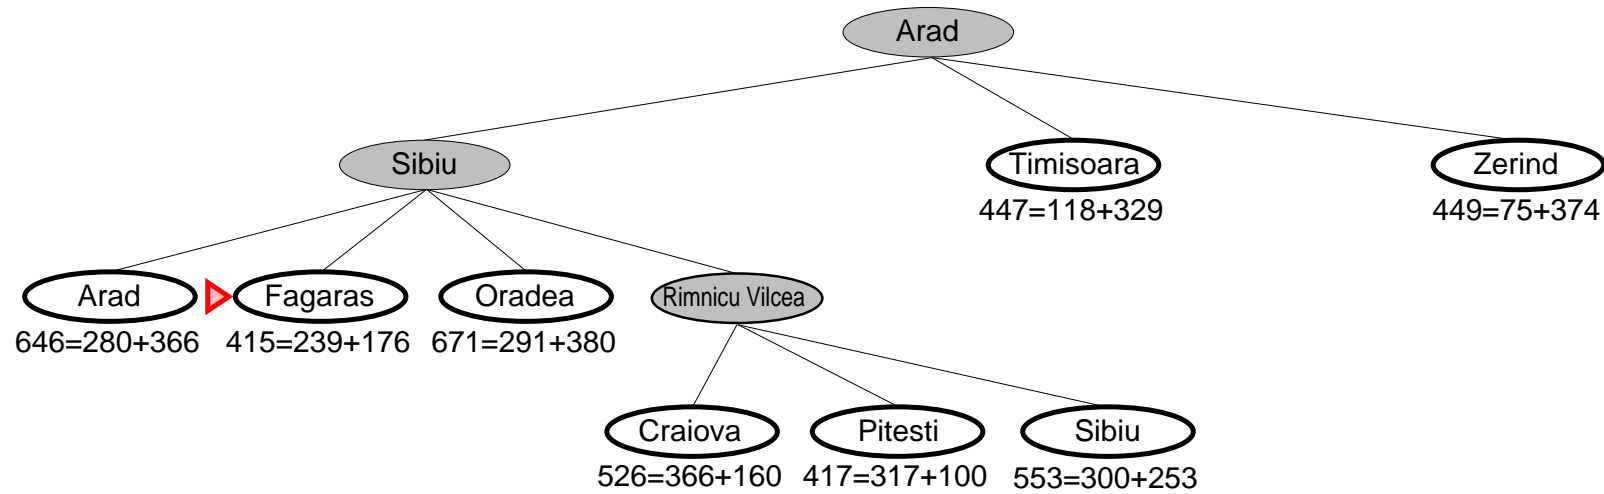


# A\* tree search example

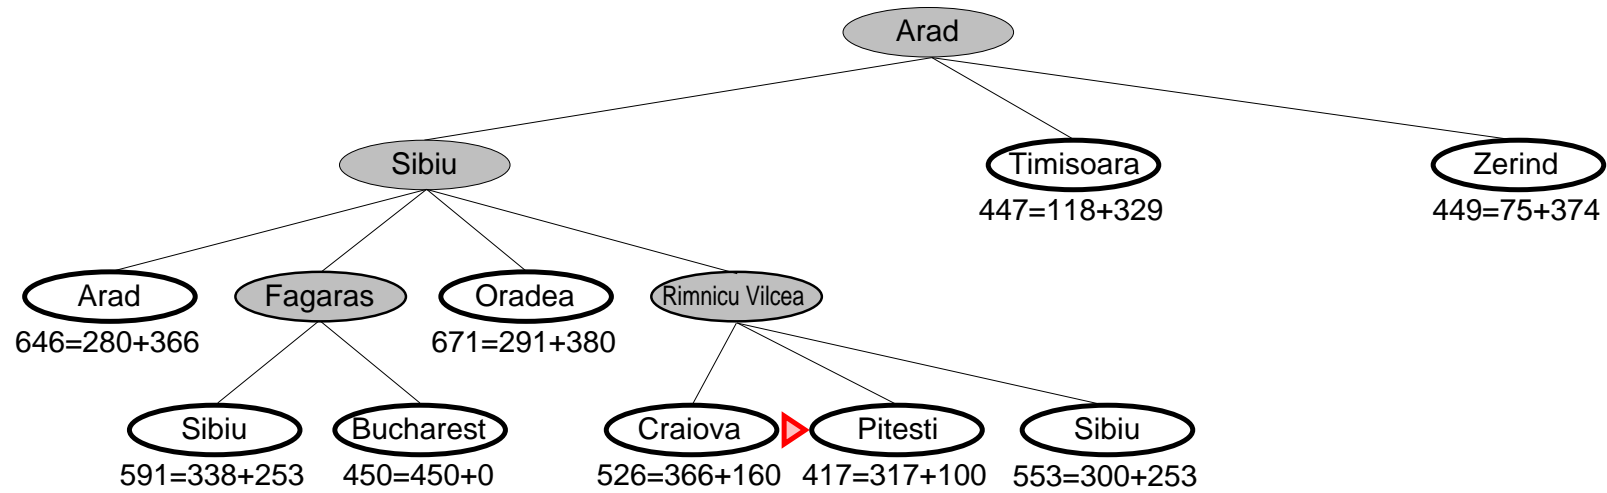




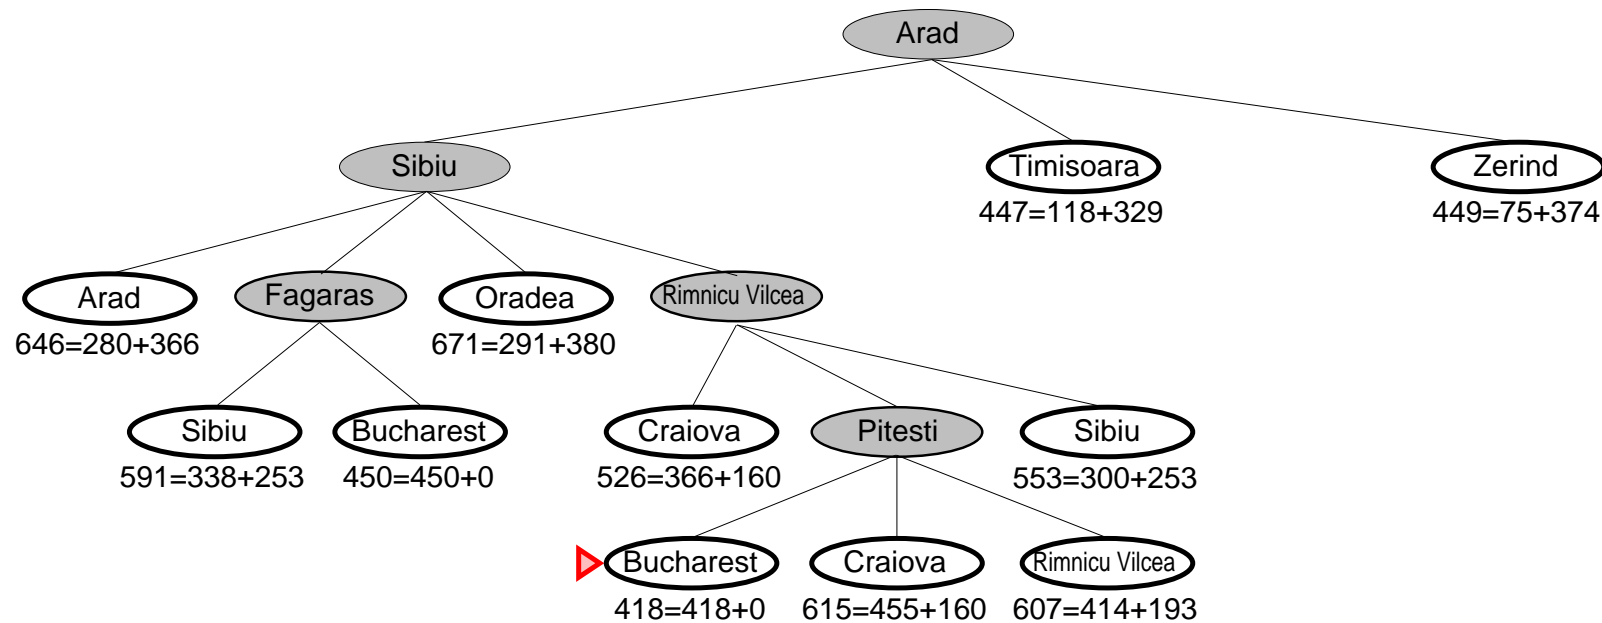
# A\* tree search example



# A\* tree search example




# A\* tree search example



# Optimality of A\* tree search

**Theorem:** A\* tree search with an admissible heuristic is optimal

**Proof.** Assume that a suboptimal goal node  $G_2$  is in the queue, and consider an arbitrary unexpanded node  $n$  on a shortest path to an optimal goal  $G$ .



$$\begin{aligned} f(G_2) &= g(G_2) && \text{since } h(G_2) = 0 \\ &> g(G) && \text{since } G_2 \text{ is suboptimal} \\ &\geq f(n) && \text{since } h \text{ is admissible} \end{aligned}$$

Since  $f(G_2) > f(n)$ , node  $G_2$  will not be selected for expansion before  $n$ .

## Properties of A\* tree search

Complete?? Yes, unless there are infinitely many nodes with  $f(n) \leq f(G)$

Time?? Exponential in [relative error in  $h \times$  length of solution.]

Space?? Keeps all nodes in memory

Optimal?? Yes

A\* expands...

- all nodes with  $f(n) < C^*$
- some nodes with  $f(n) = C^*$
- no nodes with  $f(n) > C^*$

◇ Main difficulty with A\*: memory requirements

## Optimality of A\* graph search

Reminder: graph search does not visit the same state twice

A\* graph search is not optimal for an arbitrary admissible heuristic: if a suboptimal path to a node  $n$  is discovered first, the optimal path discovered later will not be considered

Solution 1: discard the more expensive path to  $n$

Extra bookkeeping is messy, but ensures optimality.

Solution 2: ensure that optimal paths are explored first, e.g., by using a **consistent** heuristic

Enforces a form of triangle inequality (stipulating that each side of a triangle cannot be longer than the sum of the other two).

## Consistent heuristic

A heuristic is **consistent** if

$$h(n) \leq c(n, a, n') + h(n')$$


If  $h$  is consistent, then

$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n, a, n') + h(n') \\ &\geq g(n) + h(n) \\ &= f(n) \end{aligned}$$

$\Rightarrow$  I.e., along each path,  $f(n)$  is **nondecreasing**.

◇ Consistent  $\Rightarrow$  Admissible (but not the other way around)

◇ Graph search with a consistent heuristic is optimal




# Optimality of A\* graph search

**Lemma:** A\* expands nodes in order of increasing  $f$  value

Gradually adds “ $f$ -contours” of nodes (cf. breadth-first adds layers)

Contour  $i$  has all nodes with  $f = f_i$ , where  $f_i < f_{i+1}$





## Dealing with memory requirements: IDA\*

IDA\* combines iterative deepening with A\*:

- instead of depth, cutoff is the  $f$ -cost—that is, the value of  $g(n) + h(n)$

```
function IDA*(problem) returns a goal node, or failure
  root ← ROOT-NODE(problem.INITIAL-STATE)
  f_limit ←  $h(\textit{root})$ 
  loop forever
    goal, f_limit ← IDA-DFS(problem, root, f_limit)
    if goal ≠ failure return goal
    if f_limit =  $\infty$  return failure
```

## Dealing with memory requirements: IDA\*

```
function IDA-DFS(problem, node, f_limit) returns a goal node, or next cost
  f_node  $\leftarrow g(\textit{node}) + h(\textit{node})$ 
  if f_node > f_limit return (failure, f_node)
  if problem.GOAL-TEST succeeds on node.STATE return (node, f_limit)
  next_f_limit  $\leftarrow \infty$ 
  for each action in problem.ACTIONS(node.STATE) do
    successor  $\leftarrow$  CHILD-NODE(problem, node, action)
    goal, new_f_limit  $\leftarrow$  IDA-DFS(problem, successor, f_limit)
    if goal  $\neq$  failure return (goal, new_f_limit)
    next_f_limit  $\leftarrow$  min(next_f_limit, new_f_limit)
  return (failure, next_f_limit)
```

Avoids the overhead of keeping a sorted queue of nodes

Progress can be slow if  $f$ -cost increases slowly

Alternatives: recursive best-first search; (simplified) memory-bounded A\*

## Deriving heuristic functions

For example, for the 8-puzzle:

$h_1(n)$  = the number of misplaced tiles

$h_2(n)$  = the total **Manhattan** distance

- i.e., the sum over all tiles of the number of moves needed to bring the tile into required position

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

$h_1(S) = ??$  6

$h_2(S) = ??$   $4+0+3+3+1+0+2+1 = 14$

# Dominance

If  $h_2(n) \geq h_1(n)$  for all  $n$  and both heuristics are admissible, then  $h_2$  dominates  $h_1$  and is better for search

- $h_2$  is then closer to the actual cost  $h^*(n)$

Typical search costs (for 8-puzzle):

$d = 14$  IDS = 3,473,941 nodes

$A^*(h_1) = 539$  nodes

$A^*(h_2) = 113$  nodes

$d = 24$  IDS  $\approx$  54,000,000,000 nodes

$A^*(h_1) = 39,135$  nodes

$A^*(h_2) = 1,641$  nodes

Combining heuristics: given two admissible heuristics  $h_a$  and  $h_b$ ,

$$h(n) = \max(h_a(n), h_b(n))$$

is also admissible, and it dominates both  $h_a$  and  $h_b$

## Relaxed problems

Admissible heuristics can be derived from the **exact** solution cost of a **relaxed** version of the problem

If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then  $h_1(n)$  is the cost of the shortest solution


If the rules are relaxed so that a tile can move to **any adjacent square**, then  $h_2(n)$  is the cost of the shortest solution

Key point: the optimal solution cost of a relaxed problem is no greater than the optimal solution cost of the real problem

## Relaxed problems: Example

Well-known example: [traveling salesman problem](#) (TSP)

Find the shortest tour visiting all cities exactly once



[Minimum spanning tree](#) can be computed in  $O(n^2)$   
and is a lower bound on the shortest (open) tour

## Subproblems and pattern databases

Exact cost of a solution to a subproblem gives an admissible heuristic

- e.g., in an 8-puzzle,  $h(n)$  might be the **exact cost** of moving tiles 1–4 into correct position (we disregard other tiles)

Subproblems: special kinds of relaxed problems

Exact cost of subproblem solutions is often stored in a pattern database

- we precompute and store the solutions to all possible subproblems
  - can be done by searching backwards from the goal
- $h(n)$  can be obtained via simple lookup

We can add  $h$ -values for **disjoint subproblems**

## Disjoint subproblems

Subproblems of 8-puzzle:

- $h_1$ : move tiles 1–4 into appropriate position
- $h_2$ : move tiles 5–8 into appropriate position

We want to define  $h_1$  and  $h_2$  so that  $h(n) = h_1(n) + h_2(n)$  is admissible

**Not disjoint** if the cost of a subproblem is the number of total steps

- moving tiles 1–4 into position involves moving tiles 5–8 as well
- $h(n)$  is not admissible as it may count some moves twice

**Disjoint** if in each subproblem we count only the moves of the target tiles

- in  $h_1(n)$  we count only the number of moves of tiles 1–4
- in  $h_2(n)$  we count only the number of moves of tiles 5–8

To reach a goal from  $n$  we clearly need at least  $h_1(n) + h_2(n)$  steps

$\Rightarrow h(n)$  is admissible



## Summary

Heuristic functions estimate costs of shortest paths

Good heuristic can dramatically reduce search cost

Greedy best-first search expands  $n$  with lowest  $h(n)$

- incomplete and not always optimal

A\* tree search expands nodes  $n$  with lowest  $g(n) + h(n)$

- complete and optimal if heuristic is admissible
- also optimally efficient (up to tie-breaks, for forward search)

A\* graph search is optimal if heuristic is consistent

Admissible heuristic can be derived as solutions to relaxed problems

Subproblems can be seen as relaxed problems; can precompute solutions in a database